HuYuDataInsight LLC

Zhaowei Cai

2024-05-14

Q3

(a)

```
library(quantmod)
```

```
## Warning: package 'quantmod' was built under R version 4.2.3

## Loading required package: xts

## Warning: package 'xts' was built under R version 4.2.3

## Loading required package: zoo

## Warning: package 'zoo' was built under R version 4.2.3

##
## Attaching package: 'zoo'

## The following objects are masked from 'package:base':
##
##     as.Date, as.Date.numeric

## Loading required package: TTR

## Warning: package 'TTR' was built under R version 4.2.3

## Registered S3 method overwritten by 'quantmod':
##   method            from
##   as.zoo.data.frame zoo
```

```
library(urca)
```

```
## Warning: package 'urca' was built under R version 4.2.3
```

```
library(forecast)
```

```
## Warning: package 'forecast' was built under R version 4.2.3
```

```
library(tseries)
```

```
## Warning: package 'tseries' was built under R version 4.2.3
```

```
library(fGarch)
```

```
## Warning: package 'fGarch' was built under R version 4.2.3
```

```
## NOTE: Packages 'fBasics', 'timeDate', and 'timeSeries' are no longer
## attached to the search() path when 'fGarch' is attached.
##
## If needed attach them yourself in your R script by e.g.,
##         require("timeSeries")
```

```
##
## Attaching package: 'fGarch'
```

```
## The following object is masked from 'package:TTR':
##
##     volatility
```

```
library(zoo)
library(tseries)
library(rugarch)
```

```
## Warning: package 'rugarch' was built under R version 4.2.3
```

```
## Loading required package: parallel
```

```
##
## Attaching package: 'rugarch'
```

```
## The following object is masked from 'package:stats':
##
##     sigma
```

```
data = read.csv('TSLA2.csv')

closing = data$Close # closing price
log_closing = log(data$Close) # log closing price
log_return = na.omit(diff(log(data$Close))) # log return
time = as.Date(data$Date, format = '%m/%d/%y')

##Check for the trend (the Augmented Dickey-Fuller (ADF) test)
summary(ur.df(log_return, type='trend', lags=20, selectlags="BIC"))
```

```
##
## ###############################################
## # Augmented Dickey-Fuller Test Unit Root Test #
## ###############################################
##
## Test regression trend
##
##
## Call:
## lm(formula = z.diff ~ z.lag.1 + 1 + tt + z.diff.lag)
##
## Residuals:
##       Min        1Q    Median        3Q       Max
## -0.201802 -0.014870 -0.000098  0.016060  0.174363
##
## Coefficients:
##               Estimate Std. Error t value Pr(>|t|)
## (Intercept)  1.166e-04  2.169e-03   0.054    0.957
## z.lag.1     -9.360e-01  4.276e-02 -21.892   <2e-16 ***
## tt           2.152e-06  3.414e-06   0.630    0.529
## z.diff.lag  -4.573e-02  3.057e-02  -1.496    0.135
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.03448 on 1066 degrees of freedom
## Multiple R-squared:  0.492,  Adjusted R-squared:  0.4905
## F-statistic: 344.1 on 3 and 1066 DF,  p-value: < 2.2e-16
##
##
## Value of test-statistic is: -21.8922 159.7626 239.636
##
## Critical values for test statistics:
##       1pct  5pct 10pct
## tau3 -3.96 -3.41 -3.12
## phi2  6.09  4.68  4.03
## phi3  8.27  6.25  5.34
```

```
# From the result, we can see that there is no drift.
# Also, there is no linear trend for this time series because the coefficient for tt is not significant


##Check for the seasonality
n = length(log_return)
acf(log_return,main="ACF of the log return",yaxt="n")
ci=qnorm(c(0.025, 0.975))/sqrt(n)
text(y=ci,par("usr")[1],labels=round(ci,4),pos=2,xpd=TRUE)
```

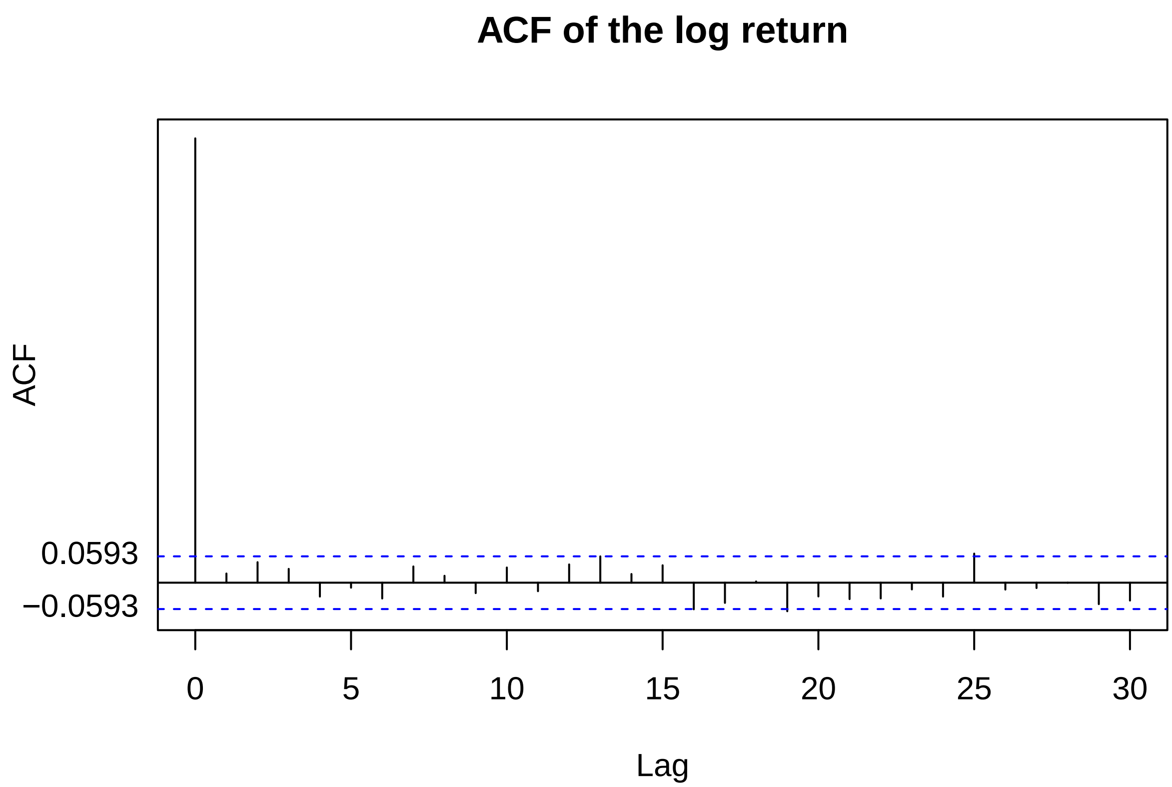


```
pacf(log_return,main="PACF of the log return",yaxt="n")
text(y=ci,par("usr")[1],labels=round(ci,4),pos=2,xpd=TRUE)
```

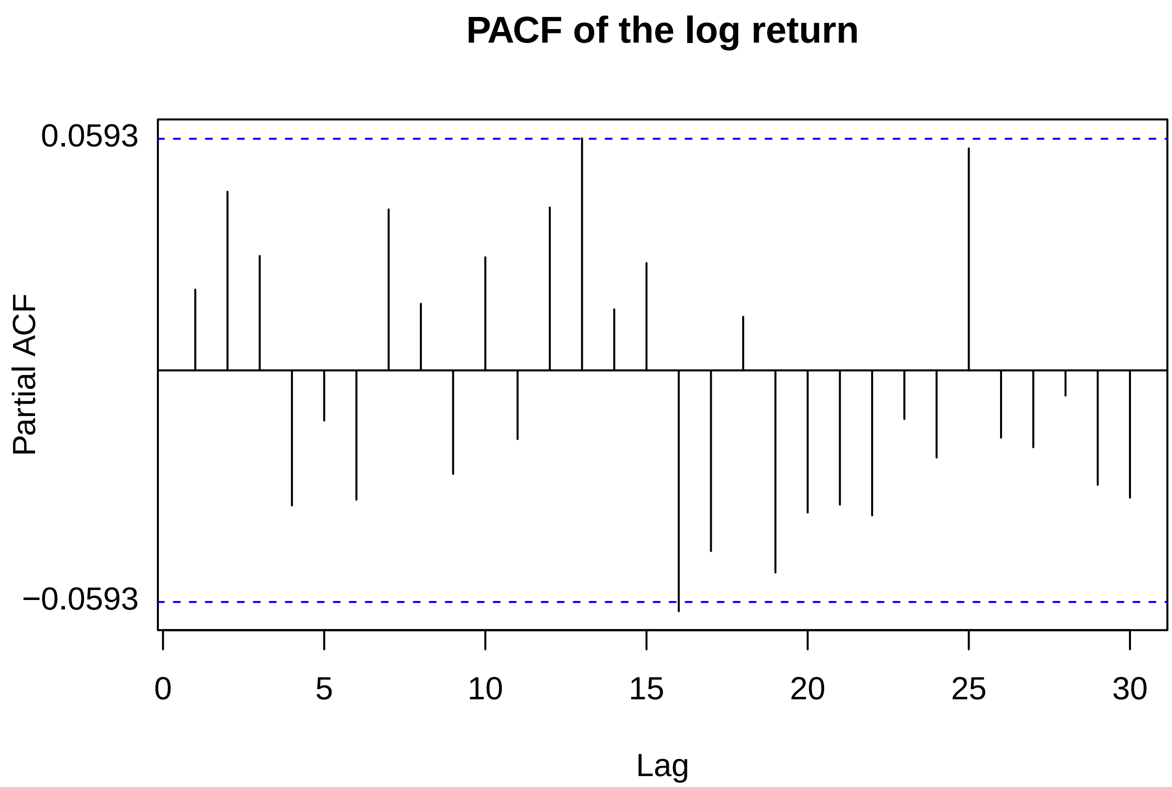


```
spec.pgram(log_return,main="Series: the log return")
```

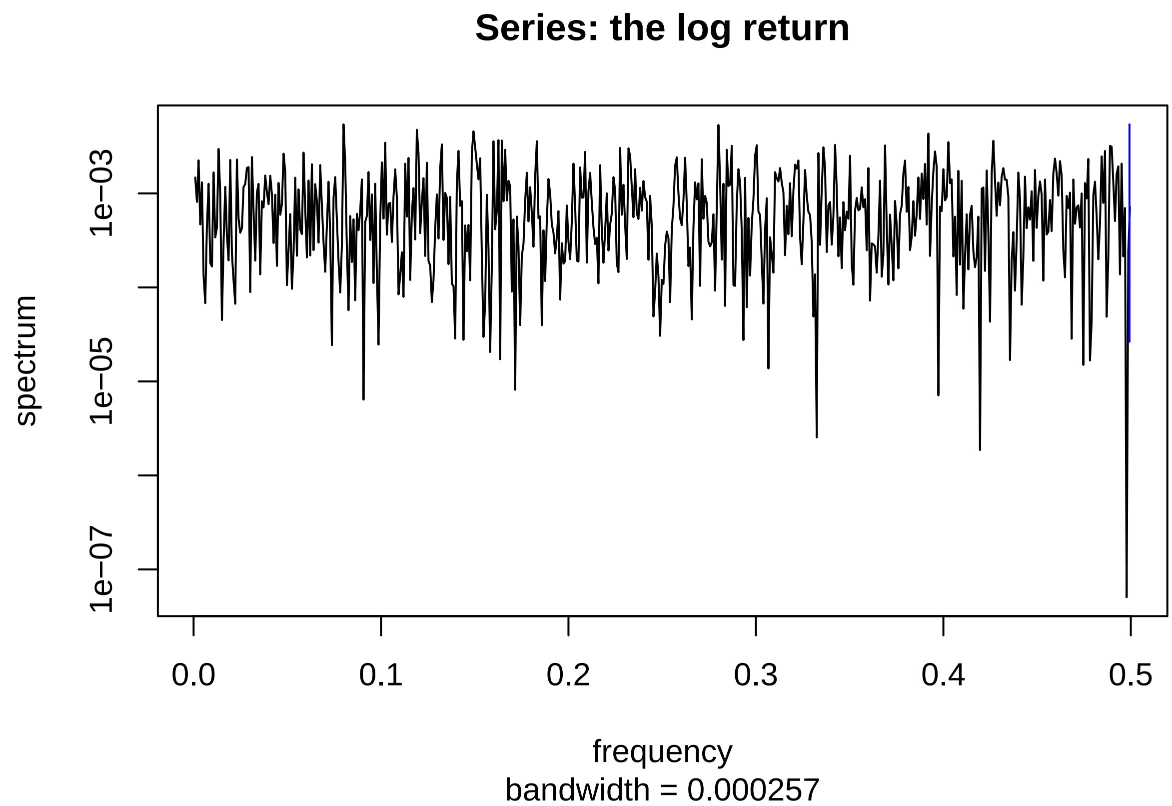


```r
# we cannot find any evidence for seasonality.

# also
adf.test(log_return)
```

```
## Warning in adf.test(log_return): p-value smaller than printed p-value
```

```
##
##  Augmented Dickey-Fuller Test
##
## data:  log_return
## Dickey-Fuller = -9.8415, Lag order = 10, p-value = 0.01
## alternative hypothesis: stationary
```

```r
# The data is stationary. Difference is not needed.
```

(b)

```r
# There is no drift or time trend
fit = auto.arima(log_return, max.p=25, max.q=25, ic="bic",
                       seasonal=F, lambda=NULL,
                       stepwise=FALSE, approximation=FALSE
                       )
summary(fit)
```

```
## Series: log_return
## ARIMA(0,0,0) with zero mean
##
## sigma^2 = 0.001182:  log likelihood = 2128.85
## AIC=-4255.69   AICc=-4255.69   BIC=-4250.7
##
## Training set error measures:
##                       ME      RMSE        MAE MPE MAPE      MASE       ACF1
## Training set 0.001132039 0.0343821 0.02288282 100  100 0.6895774 0.02069019
```

```
# ARMA(0,0)
tsdiag(fit)
```

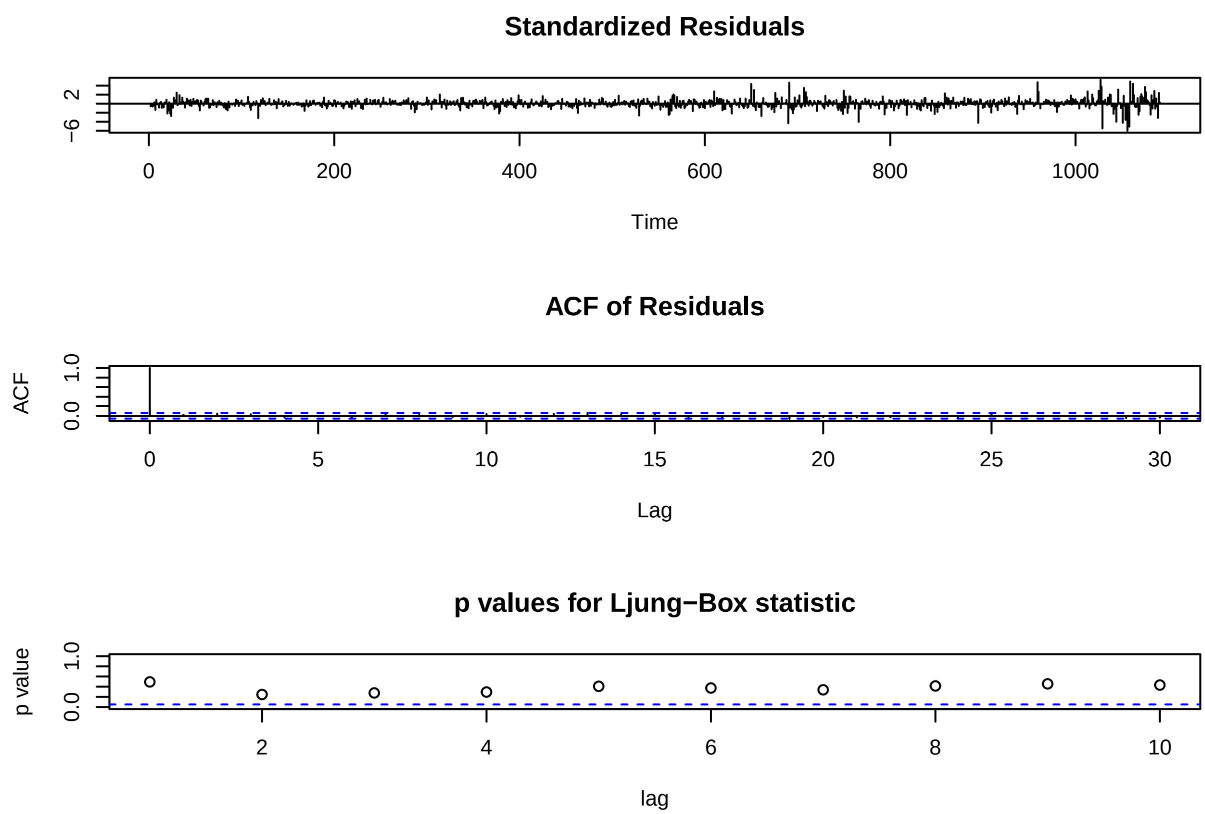


```
shapiro.test(fit$residuals)
```

```
##
##  Shapiro-Wilk normality test
##
## data:  fit$residuals
## W = 0.90385, p-value < 2.2e-16
```

```
# The null-hypothesis of this test is that the population is normally distributed.
# The null hypothesis is rejected and there is evidence that the residuals tested are not normally dist
```

(c)

```r
prediction <- forecast(fit, h=1, level=0.95) # one-day ahead log return
last_close_price <- closing[length(closing)]
(lower_interval <-as.numeric(last_close_price*exp(prediction$lower)))
```

```
## [1] 718.1479
```

```r
(price_forecast <-as.numeric(last_close_price*exp(prediction$mean)))
```

```
## [1] 768.21
```

```r
(upper_interval <-as.numeric(last_close_price*exp(prediction$upper)))
```

```
## [1] 821.762
```

```r
# Print the forecasted closing price and prediction interval
cat("1-day ahead closing price forecast:", price_forecast, "\n")
```

```
## 1-day ahead closing price forecast: 768.21
```

```r
cat("95% Prediction Interval: (", lower_interval, ", ", upper_interval, ")\n")
```

```
## 95% Prediction Interval: ( 718.1479 ,  821.762 )
```

(d)

```r
# Fit the mean model first
arma_model <- auto.arima(log_return)
arma_model
```

```
## Series: log_return
## ARIMA(0,0,0) with zero mean
##
## sigma^2 = 0.001182:  log likelihood = 2128.85
## AIC=-4255.69   AICc=-4255.69   BIC=-4250.7
```

```r
garch_spec <- ugarchspec(variance.model = list(model = "sGARCH", garchOrder = c(1,1)),
                         mean.model = list(armaOrder = c(0,0)))
garch_fit <- ugarchfit(spec = garch_spec, data = arma_model$residuals)
garch_fit
```

```
##
## *---------------------------------*
## *          GARCH Model Fit        *
## *---------------------------------*
##
## Conditional Variance Dynamics
## -----------------------------------
## GARCH Model  : sGARCH(1,1)
```

```
## Mean Model   : ARFIMA(0,0,0)
## Distribution : norm
##
## Optimal Parameters
## ------------------------------------
##         Estimate  Std. Error   t value Pr(>|t|)
## mu      0.000682    0.000862   0.79163 0.428577
## omega   0.000003    0.000004   0.80597 0.420261
## alpha1  0.028864    0.006461   4.46741 0.000008
## beta1   0.970136    0.007683 126.26416 0.000000
##
## Robust Standard Errors:
##         Estimate  Std. Error  t value Pr(>|t|)
## mu      0.000682    0.000935  0.72954  0.46567
## omega   0.000003    0.000016  0.21749  0.82783
## alpha1  0.028864    0.021416  1.34779  0.17773
## beta1   0.970136    0.026212 37.01180  0.00000
##
## LogLikelihood : 2249.554
##
## Information Criteria
## ------------------------------------
##
## Akaike       -4.1165
## Bayes        -4.0982
## Shibata      -4.1165
## Hannan-Quinn -4.1096
##
## Weighted Ljung-Box Test on Standardized Residuals
## ------------------------------------
##                         statistic p-value
## Lag[1]                      1.120  0.2899
## Lag[2*(p+q)+(p+q)-1][2]     1.412  0.3820
## Lag[4*(p+q)+(p+q)-1][5]     1.705  0.6896
## d.o.f=0
## H0 : No serial correlation
##
## Weighted Ljung-Box Test on Standardized Squared Residuals
## ------------------------------------
##                         statistic p-value
## Lag[1]                      5.099 0.02394
## Lag[2*(p+q)+(p+q)-1][5]     9.845 0.01009
## Lag[4*(p+q)+(p+q)-1][9]    11.666 0.02159
## d.o.f=2
##
## Weighted ARCH LM Tests
## ------------------------------------
##             Statistic Shape Scale P-Value
## ARCH Lag[3]     1.140 0.500 2.000  0.2857
## ARCH Lag[5]     1.346 1.440 1.667  0.6336
## ARCH Lag[7]     2.661 2.315 1.543  0.5798
##
## Nyblom stability test
## ------------------------------------
```

```
## Joint Statistic:  10.9868
## Individual Statistics:
## mu     0.06385
## omega  1.85934
## alpha1 0.46376
## beta1  0.39080
##
## Asymptotic Critical Values (10% 5% 1%)
## Joint Statistic:          1.07 1.24 1.6
## Individual Statistic:     0.35 0.47 0.75
##
## Sign Bias Test
## ------------------------------------
##                    t-value   prob sig
## Sign Bias           0.1943 0.8460
## Negative Sign Bias  1.3061 0.1918
## Positive Sign Bias  0.1315 0.8954
## Joint Effect        1.9298 0.5871
##
##
## Adjusted Pearson Goodness-of-Fit Test:
## ------------------------------------
##   group statistic p-value(g-1)
## 1    20     87.53    9.046e-11
## 2    30     99.35    1.244e-09
## 3    40    116.08    1.387e-09
## 4    50    127.56    6.369e-09
##
##
## Elapsed time : 0.1304209
```

```
forecasted_returns <- ugarchforecast(garch_fit, n.ahead = 1)
last_close_price <- closing[length(closing)]
(price_forecast <-as.numeric(last_close_price*exp(forecasted_returns@forecast$seriesFor)))
```

```
## [1] 768.7341
```

```
(lower_interval <- as.numeric(price_forecast*exp(qnorm(0.025)*forecasted_returns@forecast$sigmaFor)))
```

```
## [1] 671.6846
```

```
(upper_interval <- as.numeric(price_forecast*exp(qnorm(0.975)*forecasted_returns@forecast$sigmaFor)))
```

```
## [1] 879.806
```

```
# Print the forecasted closing price and prediction interval
cat("1-day ahead closing price forecast:", price_forecast, "\n")
```

```
## 1-day ahead closing price forecast: 768.7341
```

```r
cat("95% Prediction Interval: (", lower_interval, ", ", upper_interval, ")\n")
```

```
## 95% Prediction Interval: ( 671.6846 ,  879.806 )
```